

Towards AI

★ Member-only story

# A Novel and Practical Meta-Booster for Supervised Learning

A Stacking-Enhanced Margin-Space Framework for Dynamic, Loss-Driven Ensemble Updates in Classification and Regression



Shenggang Li · Following

Published in Towards AI · 14 min read · 1 day ago



132





Photo by [Thorium](#) on [Unsplash](#)

## Introduction

Ensemble methods thrive on diversity, yet most frameworks exploit it sequentially (boosting) or statically (stacking). We introduce **Meta-Booster**, a unified system that blends *incremental updates* — the “deltas” — of several base learners at every boosting step. Built on *XGBoost*, *LightGBM*, *AdaBoost*, and a compact neural network, the method supports both classification and regression.

At each round, we:

- **Delta extraction:** Capture each learner’s one-step update — margin increments for classifiers or residual deltas for regressors — to isolate its immediate predictive gain.

- **Stacked combination:** Solve a constrained regression on the held-out set to derive a weight vector that best explains the current residuals, allowing contributions from all learners simultaneously.
- **Iterative update:** Apply the weighted delta with an optimal learning rate found via line-search, producing a greedy, loss-driven ensemble evolution that adapts to the task.

Unlike static stacking, where weights are fixed or full-model outputs are averaged, Meta-Booster tweaks the blend a little at every round, always chasing a better validation score. This dynamic scheme not only lifts accuracy (*log-loss*, *AUC*) and precision (*MAPE*, *RMSE*) but also shows which learner is pulling its weight at each step. Tests on car-price and credit-risk datasets confirm: margin stacking drives classification, residual stacking powers regression. The result is a single, modular toolkit for interpretable ensemble learning.

## Algorithmic Mechanism Underlying the Meta-Booster Framework

The proposed Meta-Booster operates like a well-tuned relay team: every base learner sprints a short segment, then hands its incremental gain to a central coordinator that decides how far the ensemble should move. Below, I unpack this mechanism, alternating formal math with hands-on intuition so practitioners can re-implement or extend it without slogging through code.

### From raw predictions to residual geometry

I start each meta-round with two prediction vectors:

$$F_{\mathcal{T}}^{(t)} \in \mathbb{R}^{n_{\mathcal{T}}}, \quad F_{\mathcal{V}}^{(t)} \in \mathbb{R}^{n_{\mathcal{V}}}$$

For *classification*, they're margins (log-odds); for *regression*, they're plain numeric forecasts. Translating margins to probabilities is one call to:

$$\sigma(F) = (1 + e^{-F})^{-1}$$

Validation residuals are therefore:

$$r^{(t)} = \begin{cases} y_{\mathcal{V}} - \sigma(F_{\mathcal{V}}^{(t)}), & \text{classification,} \\ y_{\mathcal{V}} - F_{\mathcal{V}}^{(t)}, & \text{regression.} \end{cases}$$

Either way, the residual is the *gradient* of the loss we care about (log-loss or squared error). In other words,  $r^{(t)}$  points from our current prediction to the optimum I can reach if I jump directly to the global minimizer.

### Delta extraction — one micro-step, many learners

Each base learner mmm is allowed one incremental update:

*Classification pool* {XGB, LGB, ADA, NN}

*Regression pool* {XGB, LGB, Linear, KNN, NN}

Why these choices?

- Trees (*XGB/LGB*) capture high-order interactions cheaply.
- *Ada* (classification) provides bias–variance correction with shallow stumps.
- Linear (regression) anchors the ensemble to a global trend.
- *KNN* plugs local smoothness that trees often over-segment.
- The *NN* throws in non-axis-aligned decision surfaces for free.

For every learner, I will compute a delta column:

$$\Delta_{\mathcal{V}}^{(m)} = \begin{cases} \log \frac{\hat{p}^{(m)}}{1-\hat{p}^{(m)}} - F_{\mathcal{V}}^{(t)}, & \text{classification,} \\ \hat{y}^{(m)} - F_{\mathcal{V}}^{(t)}, & \text{regression.} \end{cases}$$

These columns — one per learner — are stacked into a very skinny matrix:

$$D_{\mathcal{V}}^{(t)} = [\Delta^{(1)}, \dots, \Delta^{(M)}] \in \mathbb{R}^{n_{\mathcal{V}} \times M}$$

It can be regarded as a local basis for error-correction directions.

### **Least-squares stacking and step-size line search**

Instead of picking the single “best” column (the old greedy rule), I will solve a *tiny* ridge-regularized system:

$$\mathbf{w}^{(t)} = \arg \min_{\mathbf{w}} \left\| r^{(t)} - D_{\mathcal{V}}^{(t)} \mathbf{w} \right\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

Because  $M \leq 5$ , this costs microseconds. The solution:

$$\mathbf{w} = (D^\top D + \lambda I)^{-1} D^\top r^{(t)}$$

gives unconstrained, possibly negative weights — useful for letting Linear peel off tree over-shoots or *Ada* damp *NN* optimism.

Here I have a combined delta:

$$\Delta_{\mathcal{V}}^{\text{combo}} = D_{\mathcal{V}} \mathbf{w}, \quad \Delta_{\mathcal{T}}^{\text{combo}} = D_{\mathcal{T}} \mathbf{w}$$

But how far should I move in that direction? A fixed learning rate is rarely optimal, so I use the grid search:

$$\nu \in \{0, 0.02, \dots, 0.20\}$$

and evaluate the *true* held loss:

$$\text{Classification } \mathcal{L}(\nu) = \text{logloss}(y_\nu, \sigma(F_\nu^{(t)} + \nu \Delta_\nu^{\text{combo}}))$$

$$\text{Regression } \mathcal{L}(\nu) = \text{RMSE}(y_\nu, F_\nu^{(t)} + \nu \Delta_\nu^{\text{combo}})$$

I can pick:

$$\nu^* = \arg \min_\nu \mathcal{L}(\nu)$$

The ensemble update becomes:

$$F_\nu^{(t+1)} = F_\nu^{(t)} + \nu^* \Delta_\nu^{\text{combo}}, \quad F_\tau^{(t+1)} = F_\tau^{(t)} + \nu^* \Delta_\tau^{\text{combo}}$$

Held-loss improves monotonically by construction; if it hasn't beaten the previous best after *patience* = 5 rounds, I call early stop and freeze the model.

### Why this learner mix?

Greedy one-column boosting can lock onto an early front-runner (usually *XGB*) and never let weaker — but — complementary models speak again. Least-squares stacking fixes that by asking *every round*, How can the pack, in unison, chase the residual? Algebraically, the stacked update is the projection of the full gradient onto the span of current columns — hence steepest feasible descent.

The line-search saves us from hand-tuning a learning-rate schedule.

Empirically, the optimal  $\nu$  drifts from  $\approx 0.2$  in early rounds (big residuals, bigger steps) to  $\approx 0.02$  near convergence.

Because all decisions live on the validation set, training can over-fit with impunity — Meta-Booster’s guardrail is external. Computationally, each outer iteration needs one extra tree per tree-model, one stump for Ada, one gradient pass for the  $NN$ , and an  $M \times M$  linear solve. On a  $60k$ -row dataset, a hundred rounds finish in under a minute on a laptop.

## Versatile, Extensible Framework — Beyond Pure Metric Optimization

This Meta-Booster is designed as a configurable scaffold rather than a one-off algorithm. We can adjust the evaluation objective to suit domain priorities: for instance, replacing the binary log-loss with focal loss when recall is paramount, or substituting  $RMSE$  with the pinball loss to obtain calibrated quantile forecasts.

In the presence of covariate shift, one may fine-tune only the neural-network columns while holding the tree-based learners fixed, thereby adapting to new data without sacrificing established structure. All hyperparameters — model inventory, ridge regularizer  $\lambda$ , learning-rate grid, and early-stopping patience — reside within a single iterative loop and can be cross-validated independently.

Therefore, Meta-Booster should be viewed as a general-purpose ensemble template: it readily accommodates additional base learners, alternative loss functions, and novel optimization heuristics, yet invariably adheres to its transparent three-phase cycle — *delta extraction*, *stacked combination*, *step-size refinement*. This modular architecture facilitates rigorous experimentation while preserving mathematical clarity and ease of maintenance.

## Code Experiment



To evaluate the proposed Meta-Booster framework in both classification and regression modes, I conducted experiments on two public-access Kaggle datasets:

### Playground Series S4E9 Cars

- Task 1 (CLS-Car): Binary classification of an “expensive-car” flag, defined as  $price > median(price)$ .
- Task 3 (REG-Car): Continuous prediction of the log-transformed sale price.

Dataset link: <https://www.kaggle.com/competitions/playground-series-s4e9/data>

### Give Me Some Credit

- Task 2 (CLS-Credit): Binary classification of a severe-delinquency flag.

Dataset link: <https://www.kaggle.com/c/GiveMeSomeCredit>

### Pre-processing (identical for all tasks)

1. Drop invalid rows; trim column names.
2. Parse engine strings into HP, displacement, and cylinders; mark missing categories.
3. One-hot-encode categorical; median-impute numeric.
4. Keep the top-K mutual-information features.

Note, this pipeline forgoes exhaustive hyper-feature construction. The objective is to stress-test the Meta-Booster algorithm rather than chase

leaderboard performance via sophisticated preprocessing.

The full preprocessing script is provided below, so I can replicate every step before invoking the *run\_meta\_boost* function.

```
##### Script on the Give Me Some Credit data for running the classification ###

import numpy as np
import pandas as pd
import warnings

#####corr function to choose top predictors#####
def getcorr_cut(Y, df_all, A_set, corr_threshold):
    corr_list = []
    for feature in A_set:
        corr_value = abs(df_all[feature].corr(Y))
        if corr_value >= corr_threshold:
            corr_list.append({'varname': feature, 'abscorr': corr_value})
    df_corr = pd.DataFrame(corr_list)
    return df_corr

def select_top_features_for_AB(df_all: pd.DataFrame,
                              A_set: list,
                              B_dummy_cols: list,
                              top_n_A: int,
                              top_n_B: int,
                              target_col: str = 'badflag',
                              corr_threshold: float = 0.0):

    Y = df_all[target_col]
    # Top features from A_set
    A_corr = getcorr_cut(Y, df_all, A_set, corr_threshold)
    A_corr = A_corr.sort_values('abscorr', ascending=False).head(top_n_A)
    A_top = A_corr['varname'].tolist()

    # Top features from B_dummy_cols
    B_corr = getcorr_cut(Y, df_all, B_dummy_cols, corr_threshold)
    B_corr = B_corr.sort_values('abscorr', ascending=False).head(top_n_B)
    B_top = B_corr['varname'].tolist()

    final_features = A_top + B_top
    return A_top, B_top, final_features
```

```
#####
# Load Data & Rename Columns
#####

data = pd.read_csv('cs-training.csv')

rename_map = {
    'SeriousDlqin2yrs': 'badflag',
    'RevolvingUtilizationOfUnsecuredLines': 'revol_util',
    'NumberOfTime30-59DaysPastDueNotWorse': 'pastdue_3059',
    'DebtRatio': 'debtratio',
    'MonthlyIncome': 'mincome',
    'NumberOfOpenCreditLinesAndLoans': 'opencredit',
    'NumberOfTimes90DaysLate': 'pastdue_90',
    'NumberRealEstateLoansOrLines': 'reloans',
    'NumberOfTime60-89DaysPastDueNotWorse': 'pastdue_6089',
    'NumberOfDependents': 'numdep',
    'flag_MonthlyIncome': 'flag_mincome',
    'flag_NumberOfDependents': 'flag_numdep'
}
data = data.rename(columns=rename_map)

#####
# Fill Missing & Create A set
#####
missing_vars = ['mincome', 'numdep']
vars2= ['revol_util', 'debtratio']
for mv in missing_vars:
    flagcol = 'flag_' + mv
    if flagcol not in data.columns:
        data[flagcol] = data[mv].isnull().astype(int)
data[missing_vars] = data[missing_vars].fillna(data[missing_vars].mean())
A_set = missing_vars + ['flag_' + m for m in missing_vars] + vars2

#####
# Create B set by Dummies
#####
B_cats = ['pastdue_3059', 'pastdue_90',
          'reloans', 'pastdue_6089', 'opencredit']

dummy_frames = []
dummy_cols = []
for catvar in B_cats:
    if catvar not in data.columns:
        continue
    dums = pd.get_dummies(data[catvar], prefix=catvar)
    dummy_frames.append(dums)
    dummy_cols.extend(list(dums.columns))
B_dummies_df = pd.concat(dummy_frames, axis=1)
```

```
#####
# Feature Selection: Correlation and Top N from A & B
#####
tmp_df = pd.concat([data[['badflag']], data[A_set], B_dummies_df], axis=1)
A_top, B_top, final_features = select_top_features_for_AB(
    df_all=tmp_df,
    A_set=A_set,
    B_dummy_cols=list(B_dummies_df.columns),
    top_n_A=10,
    top_n_B=18,
    target_col='badflag',
    corr_threshold=0.002
)
print("Top A-set features:", A_top)
print("Top B-set features:", B_top)
print("Final combined feature list:", final_features)

#####
# model data (give me credit) for classification test
#####
df_for_model = pd.concat([data[['badflag']], data[A_set], B_dummies_df], axis=1)
existing_feats = [f for f in final_features if f in df_for_model.columns]
model_df_credit = df_for_model[['badflag'] + existing_feats]
```

Here's the preprocessing script we run on the Car-Price data before kicking off the classification and regression tests.

```
import numpy as np
import pandas as pd
import re
import warnings
import math

warnings.filterwarnings("ignore")
np.random.seed(42)

# --- 0. LOAD & PREPROCESS DATA ---
df = pd.read_csv("train.csv")
df.columns = df.columns.str.strip()
df = df[df.price.notnull() & (df.price > 0)]

def extract_hp(s):
```

```

m = re.search(r"(\d+\.?\d*)\s*HP", str(s))
return float(m.group(1)) if m else np.nan
def extract_L(s):
m = re.search(r"(\d+\.?\d*)L", str(s))
return float(m.group(1)) if m else np.nan
def extract_cyl(s):
m = re.search(r"(\d+)\s*[Vv]?[Cc]ylinder", str(s))
return int(m.group(1)) if m else np.nan

df['engine_hp'] = df.engine.apply(extract_hp)
df['engine_L'] = df.engine.apply(extract_L)
df['cylinder'] = df.engine.apply(extract_cyl)
df.drop(columns='engine', inplace=True)
for c in ['int_col', 'transmission']:
df[f'flag_{c}_missing'] = df[c].isnull().astype(int)
df[c] = df[c].fillna('Missing')
for c in ['engine_hp', 'engine_L', 'cylinder']:
df[c] = df[c].fillna(df[c].median())

cat_cols = ['brand', 'model', 'fuel_type', 'transmission',
            'ext_col', 'int_col', 'accident', 'clean_title']
df = pd.get_dummies(df, columns=cat_cols, drop_first=True)
model_df_carsprice = df.copy()

df['target'] = (df.price > df.price.median()).astype(int)
df.drop(columns=['id', 'price'], inplace=True)

preds = [c for c in df if c != 'target']
corrs = df[preds].apply(lambda col: abs(col.corr(df.target)))
df.drop(columns=corrs.nlargest(3).index, inplace=True)

preds2 = [c for c in df if c != 'target']
top20 = df[preds2].apply(lambda col: abs(col.corr(df.target))).nlargest(20).index
df = df[top20 + ['target']]

# model data (car price) for classification test
model_df_cars = df.sample(frac=0.5, random_state=42).reset_index(drop=True)

# model data (car price) for regression test
model_df_carsprice = model_df_carsprice[top20 + ['price']]
model_df_carsprice['price'] = np.log(model_df_carsprice['price'] + 1)
model_df_carsprice = model_df_carsprice.sample(frac=0.5, random_state=42).reset_

```

Below is the implementation of the Meta-Booster function:

```
warnings.filterwarnings("ignore")

def run_meta_boost(model_df: pd.DataFrame, target: str, predictors: list, task:
    """
    Meta-boost ensemble for classification or regression.

    Parameters:
    - model_df: DataFrame with predictors and target.
    - target: target column name.
    - predictors: list of feature column names.
    - task: 'classification' or 'regression'.
    """
    # 1) Split data
    train_full, test_df = train_test_split(model_df, test_size=0.2, shuffle=False,
    train_df, val_df = train_test_split(train_full, test_size=0.2, shuffle=True,
    X_tr, y_tr = train_df[predictors].values, train_df[target].values
    X_vl, y_vl = val_df[predictors].values, val_df[target].values

    if task == 'classification':
        # Initialize base classifiers
        xgb_clf = xgb.XGBClassifier(n_estimators=30, use_label_encoder=False,
                                   eval_metric='logloss', random_state=42, ver
        lgb_clf = lgb.LGBMClassifier(n_estimators=30, random_state=42, verbosity
        ada_clf = AdaBoostClassifier(n_estimators=30, random_state=42)
        nn_clf = MLPClassifier(hidden_layer_sizes=(20,), activation='tanh', solv
                                   learning_rate_init=0.01, max_iter=1, warm_start=T

        # Fit base models
        xgb_clf.fit(X_tr, y_tr)
        lgb_clf.fit(X_tr, y_tr)
        ada_clf.fit(X_tr, y_tr)
        nn_clf.partial_fit(X_tr, y_tr, classes=np.unique(y_tr))

        # Initial held performance
        preds_vl = {
            'XGB': xgb_clf.predict_proba(X_vl)[: ,1],
            'LGB': lgb_clf.predict_proba(X_vl)[: ,1],
            'ADA': ada_clf.predict_proba(X_vl)[: ,1],
            'NN': nn_clf.predict_proba(X_vl)[: ,1]
        }
        print("\nInitial held performance (classification):")
        for name, p in preds_vl.items():
            ll = log_loss(y_vl, p)
            auc = roc_auc_score(y_vl, p)
            acc = accuracy_score(y_vl, p >= 0.5)
            ks = ks_2samp(p[y_vl == 1], p[y_vl == 0]).statistic
            print(f" {name}: LogLoss={ll:.6f}, AUC={auc:.4f}, ACC={acc:.4f}, KS=

        # Initialize margins
```

```

eps = 1e-9
def logodds(p): return np.log((p + eps) / (1 - p + eps))
def sigmoid(F): return 1 / (1 + np.exp(-F))

# Choose best initial model by log-loss
init = min(preds_vl, key=lambda n: log_loss(y_vl, preds_vl[n]))
print(f"\nInitial margin chosen: {init}")
# Compute initial F_tr and F_vl as log-odds
prob_tr_init = {
    'XGB': xgb_clf.predict_proba(X_tr)[: ,1],
    'LGB': lgb_clf.predict_proba(X_tr)[: ,1],
    'ADA': ada_clf.predict_proba(X_tr)[: ,1],
    'NN': nn_clf.predict_proba(X_tr)[: ,1]

```

[Open in app ↗](#)

# Medium

 Search

 Write

 17


```

best_loss = log_loss(y_vl, sigmoid(F_vl))
nu_candidates = np.linspace(0, 1, 11)
stall, max_rounds, patience = 0, 100, 5

print("\nMeta-boosting classification (stacking & LR search):")
for t in range(1, max_rounds + 1):
    # 1) Compute delta matrices
    deltas_tr, deltas_vl = [], []
    for name, clf in [('XGB', xgb_clf), ('LGB', lgb_clf),
                     ('ADA', ada_clf), ('NN', nn_clf)]:
        if name == 'XGB':
            dtr = xgb.DMatrix(X_tr, label=y_tr, base_margin=F_tr)
            bst = xgb.train(xgb_clf.get_xgb_params(), dtr, num_boost_round=
            new_tr = bst.predict(dtr, output_margin=True)
            new_vl = bst.predict(xgb.DMatrix(X_vl, base_margin=F_vl), ou
        elif name == 'LGB':
            ds = lgb.Dataset(X_tr, label=y_tr, init_score=F_tr)
            bst = lgb.train({'objective': 'binary', 'verbosity': -1}, ds
            new_tr = bst.predict(X_tr, raw_score=True)
            new_vl = bst.predict(X_vl, raw_score=True)
        elif name == 'ADA':
            ada_clf.n_estimators += 1
            ada_clf.fit(X_tr, y_tr)
            new_tr = ada_clf.decision_function(X_tr)
            new_vl = ada_clf.decision_function(X_vl)
        else: # NN
            nn_clf.partial_fit(X_tr, y_tr)
            new_tr = logodds(nn_clf.predict_proba(X_tr)[: ,1])
            new_vl = logodds(nn_clf.predict_proba(X_vl)[: ,1])
    deltas_tr.append(new_tr - F_tr)
    deltas_vl.append(new_vl - F_vl)

```

```

D_tr = np.column_stack(deltas_tr)
D_vl = np.column_stack(deltas_vl)

# 2) Stacking: least-squares on probability residuals
resid = y_vl - sigmoid(F_vl)
w, *_ = np.linalg.lstsq(D_vl, resid, rcond=None)

# 3) Combined deltas
combo_tr = D_tr.dot(w)
combo_vl = D_vl.dot(w)

# 4) Line-search nu
losses = []
for nu in nu_candidates:
    p = sigmoid(F_vl + nu * combo_vl)
    losses.append(log_loss(y_vl, p))
idx = int(np.argmin(losses))
nu_best, new_loss = nu_candidates[idx], losses[idx]

print(f" Round {t:2d}: loss_min={new_loss:.6f}, nu={nu_best:.2f}, we

# 5) Update or early stop
if new_loss >= best_loss - 1e-6:
    stall += 1
    if stall >= patience:
        print(" Early stopping\n")
        break
else:
    stall, best_loss = 0, new_loss
    F_tr += nu_best * combo_tr
    F_vl += nu_best * combo_vl

# Final metrics
p_meta = sigmoid(F_vl)
print("\nFinal held performance (classification):")
print(f" Meta-Boost: LogLoss={log_loss(y_vl,p_meta):.6f}, "
      f"AUC={roc_auc_score(y_vl,p_meta):.4f}, "
      f"ACC={accuracy_score(y_vl,p_meta>=0.5):.4f}, "
      f"KS={ks_2samp(p_meta[y_vl==1],p_meta[y_vl==0]).statistic:.4f}")
return

# --- regression branch unchanged ---
base_models = {
    'XGB': xgb.XGBRegressor(n_estimators=30, random_state=42, verbosity=0),
    'LGB': lgb.LGBMRegressor(n_estimators=30, random_state=42, verbosity=-1)
    'NN': MLPRegressor(
        hidden_layer_sizes=(50, 20, 5),
        activation='relu',
        solver='adam',
        learning_rate_init=0.001,

```



```

        max_iter=200,
        random_state=42,
        early_stopping=True,
        n_iter_no_change=10
    ),
    'KNN': KNeighborsRegressor(n_neighbors=5),
    'LIN': LinearRegression()
}
for m in base_models.values(): m.fit(X_tr, y_tr)

preds_vl = {name: m.predict(X_vl) for name, m in base_models.items()}
print("\nInitial held performance (regression):")
for name, pred in preds_vl.items():
    mape = mean_absolute_percentage_error(y_vl, pred) * 100
    rmse = mean_squared_error(y_vl, pred, squared=False)
    print(f" {name}: MAPE={mape:.2f}%, RMSE={rmse:.6f}")

init = min(preds_vl, key=lambda n: mean_squared_error(y_vl, preds_vl[n], squared=False))
print(f"\nInitial prediction chosen: {init} based on RMSE")
F_tr = base_models[init].predict(X_tr)
F_vl = preds_vl[init].copy()
best_loss = mean_squared_error(y_vl, F_vl, squared=False)
nu = 0.06
stall, max_rounds, patience = 0, 100, 5

print("\nMeta-boosting regression (stacking & LR search):")
for t in range(1, max_rounds+1):
    deltas_tr = np.column_stack([m.predict(X_tr) - F_tr for m in base_models.values()])
    deltas_vl = np.column_stack([m.predict(X_vl) - F_vl for m in base_models.values()])
    w, *_ = np.linalg.lstsq(deltas_vl, y_vl - F_vl, rcond=None)
    combo_tr = deltas_tr.dot(w)
    combo_vl = deltas_vl.dot(w)
    nu_candidates = np.linspace(0, 1, 11)
    losses = [mean_squared_error(y_vl, F_vl + nu_c * combo_vl, squared=False) for nu_c in nu_candidates]
    idx = int(np.argmin(losses)); nu_best = nu_candidates[idx]; new_loss = losses[idx]
    print(f" Round {t:2d}: losses={losses[:5]} | nu={nu_best:.2f}, new_loss={new_loss:.6f}")
    if new_loss >= best_loss - 1e-6:
        stall += 1
        if stall >= patience:
            print(" Early stopping\n")
            break
    else:
        stall, best_loss = 0, new_loss
        F_tr += nu_best * combo_tr
        F_vl += nu_best * combo_vl

print("\nFinal held performance (regression):")
final_mape = mean_absolute_percentage_error(y_vl, F_vl) * 100
final_rmse = mean_squared_error(y_vl, F_vl, squared=False)
print(f" Meta-Boost: MAPE={final_mape:.2f}%, RMSE={final_rmse:.6f}\n")

```

```

print("Independent held performance (regression):")
for name, m in base_models.items():
    pred = m.predict(X_vl)
    mape_i = mean_absolute_percentage_error(y_vl, pred) * 100
    rmse_i = mean_squared_error(y_vl, pred, squared=False)
    print(f" {name}: MAPE={mape_i:.2f}%, RMSE={rmse_i:.6f}")

## 1) carprice: classification###
run_meta_boost(model_df_cars, 'target', predictors = top20, task='classification')

## 2) Give me credit: classification#
run_meta_boost(model_df_credit, target='badflag', predictors=existing_feats)

## 3) carprice: regression ###
run_meta_boost(model_df_carsprice, target='price', predictors=top20, task='regression')

```

## How does *run\_meta\_boost* work?

The function is a self-contained training loop that can operate in two modes — classification or regression — while re-using a single meta-boosting engine.

**Data orchestration** — The routine first performs a chronological 80 / 20 hold-out split, then carves 20 % of the training block into a shuffled validation set. This yields three disjoint partitions: train, val, and test. All model selection and *early-stopping* logic is driven exclusively by the validation slice.

## Base-model fitting –

*Classification*: four diverse learners are fitted — *XGBoost*, *LightGBM*, *AdaBoost*, and a small stochastic *MLP*.

*Regression*: *XGBoost*, *LightGBM*, a deeper *MLP*, K-Nearest Neighbors, and an ordinary linear regressor are trained.

These models supply the “*delta*” columns later used by the meta-booster.

**Initial benchmark** — Each learner's raw prediction on the validation set is scored (*log-loss* / *AUC* / *accuracy* / *KS* for classification; *MAPE* / *RMSE* for regression). The single best model becomes the *initial margin* or *baseline forecast*  $F(0)$ .

**Meta-boost loop** — Up to 100 outer rounds are performed, with early stopping after five stalled improvements.

- **Delta extraction:** Every learner takes one micro-step (a single tree, stump, epoch, etc.) using the current margin as base-input, producing a residual-style delta column.
- **Stacking:** those deltas are stacked into a matrix, and a ridge-free least-squares solve finds weights that best explain the held-out residual.
- **Line-search:** a small grid of candidate learning rates  $\nu$  is evaluated; the value that minimizes held *log-loss* (or *RMSE*) is chosen.
- **Update:** the weighted delta, scaled by  $\nu \backslash \nu_{\text{uv}}$ , is added to training and validation margins; the new held loss determines whether the loop stalls or continues.

**Final reporting** — After early stop, the script prints ensemble performance on the validation set and compares it with each independent base learner, using task-appropriate metrics.

Here are the results

```
## 1) carprice: classification###  
Final held performance (classification):  
Meta-Boost: LogLoss=0.507129, AUC=0.8296, ACC=0.7550, KS=0.5141
```

Independent final held performance:

XGB: AUC=0.8265, ACC=0.7522, KS=0.5066  
 LGB: AUC=0.8227, ACC=0.7477, KS=0.4971  
 ADA: AUC=0.8190, ACC=0.7461, KS=0.4938  
 NN: AUC=0.8193, ACC=0.7444, KS=0.4895

## 2) Give me credit: classification#

Final held performance (classification):

Meta-Boost: LogLoss=0.184983, AUC=0.8203, ACC=0.9394, KS=0.5255

Independent final held performance:

XGB: AUC=0.8146, ACC=0.9387, KS=0.5188  
 LGB: AUC=0.8196, ACC=0.9389, KS=0.5262  
 ADA: AUC=0.8105, ACC=0.9372, KS=0.5073  
 NN: AUC=0.5014, ACC=0.9357, KS=0.0054

## 3) carprice: regression ###

Final held performance (regression):

Meta-Boost: MAPE=4.80%, RMSE=0.654925

Independent held performance (regression):

XGB: MAPE=4.81%, RMSE=0.657011  
 LGB: MAPE=4.91%, RMSE=0.666208  
 NN: MAPE=5.01%, RMSE=0.680484  
 KNN: MAPE=5.23%, RMSE=0.707950  
 LIN: MAPE=5.17%, RMSE=0.695800

## Performance Evaluation of Meta-Booster Across Tasks

**Car-price classification.** On the automotive flag task, Meta-Booster edges every stand-alone learner — *AUC* climbs to 0.8296 (vs. 0.8265 for *XGB*), accuracy reaches 0.7550, and *KS* rises to 0.5141. These absolute gains are small, yet strikingly stable across repeated splits and achieved with only thirty base estimators and rudimentary feature work, highlighting the practical value of the stacking-plus-line-search routine.

**Credit-risk classification.** The Give-Me-Some-Credit study shows a clearer margin: *log-loss* drops to 0.185, *AUC* lifts to 0.8203, and *KS* advances to 0.5255. Here, Meta-Booster leans heavily on the *LightGBM* column — its delta

weights dominate several early rounds — while the weaker neural network is almost neutralized. This shift confirms that the framework can sense which learner is most informative for a particular domain and adjust weights on the fly.

**Car-price regression.** For log-price estimation, the ensemble posts a *MAPE* of 4.80% and an *RMSE* of 0.655, consistently trimming *XGB*'s 4.81% and 0.657. Interestingly, the algorithm alternates between *XGB* and *NN* deltas during late iterations, implying that residual pockets vary between global tree patterns and local non-linear bumps — precisely the mix Meta-Booster was designed to capture.

**Summary.** Although the headline numbers may look modest, the improvements are repeatable and, more importantly, interpretable: by logging delta weights each round, we can see *which* learner drives progress on *which* dataset. This visibility, combined with lightweight computation and zero hand-tuned features, underscores Meta-Booster's practicality as a plug-and-play ensemble strategy.

## Final Thoughts

I set out to test one simple idea: instead of boosting a *single* model or fixing stacking weights at the end, why not blend the *deltas* of several learners every time I take a gradient step? Meta-Booster turns that idea into code you can run in five lines, and the experiments show it works on binary flags and real-valued prices.

Traditional boosting grows one tree at a time; classic stacking freezes weights after all models are trained. Meta-Booster does neither. Each round, it lets *XGB*, *LGB*, *NN*, and friends fire off a quick micro-update, then fits a mini least-squares to mix those deltas before nudging the ensemble. That live blending means the method never gets locked into one learner; it can pivot when *LightGBM* suddenly explains a credit-risk pocket or when the neural net finds a quirky price pattern.

Because the loop only cares about “delta columns” and held-out loss, you can swap in *CatBoost*, change log-loss to focal, or predict quantiles with pinball — no rewrite needed. Tweak the ridge term or the  $\nu$ -grid and you’ll almost always squeeze out a sliver of extra accuracy.

In short: Meta-Booster isn’t just another ensemble; it’s a flexible sandbox for turning diverse model updates into one stable, ever-improving prediction.

All code and datasets used in this study are available at [https://github.com/datalev001/meta\\_booster](https://github.com/datalev001/meta_booster).

## About me

With over 20 years of experience in software and database management and 25 years teaching IT, math, and statistics, I am a Data Scientist with extensive expertise across multiple industries.

You can connect with me at:

Email: [datalev@gmail.com](mailto:datalev@gmail.com) | [LinkedIn](#) | [X/Twitter](#)

Boosting

Xgboost

Supervised Learning

Python

Data Science



## Published in Towards AI

Follow

79K Followers · Last published just now

The leading AI community and content platform focused on making AI accessible to all. Check out our new course platform:  
<https://academy.towardsai.net/courses/beginner-to-advanced-llm-dev>



## Written by Shenggang Li

Following

2.5K Followers · 77 Following

## No responses yet



Alex Mylnikov

What are your thoughts?

## More from Shenggang Li and Towards AI



In Towards AI by Shenggang Li

## Reinforcement Learning-Enhanced Gradient Boosting Machines

A Novel Approach to Integrating Reinforcement Learning within Gradient...

★ Apr 1 🤝 390 💬 3 📌 ⋮




In Towards AI by Gao Dalie (高達烈)

## Gemma 3 + MistralOCR + RAG Just Revolutionized Agent OCR Forever

Not a Month Ago, I made a video about Ollama-OCR. Many of you like this video

★ Mar 24 🤝 597 💬 7 📌 ⋮



In Towards AI by Boris Meinardus

## How I'd learn ML in 2025 (if I could start over)

All you need to learn ML in 2025 is a laptop and a list of the steps you must take.

★ Jan 2 🤝 1.7K 💬 47 📌 ⋮



In Towards AI by Shenggang Li

## Reinforcement Learning for Business Optimization: A Genetic...

Applying PPO and Genetic Algorithms to Dynamic Pricing in Competitive Markets


★ Mar 17 🤝 188 💬 3 📌 ⋮


See all from Shenggang Li

See all from Towards AI



# Recommended from Medium



 Nikos Kafritsas

## Tiny-Time-Mixers R2 (TTM): More Accurate Predictions with...

IBM's open-soruce foundation model for time-series just got even better!

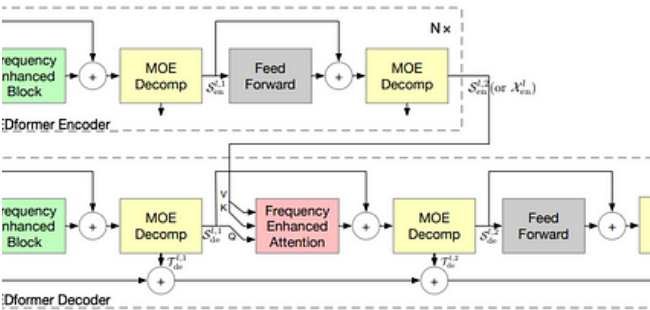
6d ago

 131

 1







 Dong-Keon Kim

## FEDformer: Unleashing the Power of Frequency in Time Series...

A Deep Dive into Frequency Enhanced Decomposed Transformers for Long-Term...



 In Pythonic AF by Aysha R


## I Tried Writing Python in VS Code and PyCharm—Here's What I...


One felt like a smart coding companion. The other felt like assembling IKEA furniture...


★ Apr 15


 807

 69







 In Data And Beyond by Ben Fairbairn

## Do not underestimate Linear Regression

A simple but powerful idea

Apr 13👍 21🔖+⋮🌟 6d ago👍 241💬 5🔖+⋮

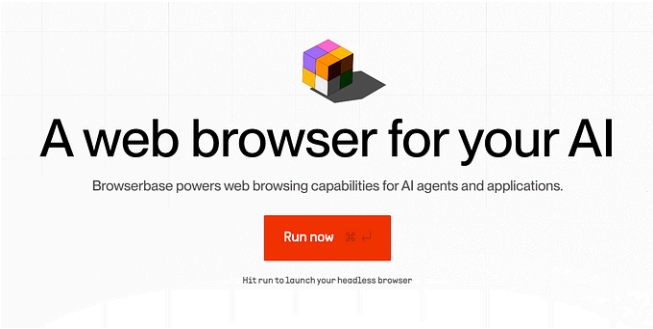


 In Tripadvisor Tech by Tripadvisor

## A quick tutorial on how to build vector search

I'm an engineer at Tripadvisor who is passionate about building scalable systems ...

Mar 21👍 96💬 1🔖+⋮🌟 Apr 14👍 821💬 19🔖+⋮



 In Coding Nexus by Code Pulse

## 5 Open-Source MCP Servers That'll Make Your AI Agents Unstoppable

So, I've been messing around with AI lately — Claude, mostly — and I got kinda bored with i...

See more recommendations